

Towards a Verified Opaque Semantics for Software Transactional Memory

RODRIGO RIBEIRO

Universidade Federal de Ouro Preto, Minas Gerais, Brazil

ANDRÉ RAUBER DU BOIS

Universidade Federal de Pelotas, Rio Grande do Sul, Brazil

(e-mail: rodrigo@decsi.ufop.br)

Abstract

Here goes the abstract...

1 Introduction

It is a well known fact that writing correct lock based concurrent programs is a painful task, since locks doesn't scale (Harris *et al.*, 2005). *Software Transactional Memory* (STM) is a promising approach to write concurrent software since it can perform groups of memory operations atomically (Shavit & Touitou, 1995) thus providing automatic rollback, deadlock, priority inversion and lock granularity freedom.

Haskell's implementation of STM pioneered, thought its type system, the strict separation between transactional and non-transactional code and its interface allows the compositional definition of transactions (Harris *et al.*, 2005).

Talk about models / semantics of STM. Mention Hutton work.

More specifically, our contributions are:

- We define a trace-based small step semantics for a high-level language with STM support.
- A Haskell implementation of opacity property for traces generated by proposed language semantics and its test using QuickCheck (Claessen & Hughes, 2000) and HPC (Gill & Runciman, 2007).

Our work shares some similarities with (?): We use a small variation of its reduced language for STM, but our focus is on verify a safety property (namely, opacity ()) for STM, instead of a compiler correctness theorem for transactional virtual machine.

2 A Model for Software Transactional Memory

In order to analyse a semantics for a high level language with transactional memory support, we follow the same approach used by (Hu & Hutton, 2009): defining a simple lan-

guage of integer-valued expressions (e.g. Figure 1) extended with conditionals, *orElse* and *retry* constructs.

We let meta-variable v denote integer constants and x (transactional) variables.

$p ::=$	v	$t ::=$	v
	$\text{fork } p$		$\text{readTVar } x$
	$\text{Atomic } t$		$\text{writeTVar } x \ t$
	$p \oplus p$		$t \oplus t$
			$\text{if } t \text{ then } t \text{ else } t$
			retry
			$t \text{ orElse } t$

Fig. 1. Language syntax.

In our semantics, we use finite maps to represent heaps and logs used by transactions (i.e. read and write sets). Notation \bullet denotes an empty finite mapping. Variable Θ represents a triple formed by a heap, read and write set used by some transaction t . We represent read and write sets and the heap by Θ_r , Θ_w and Θ_h respectively. Let $h(x)$ denotes the operation of retrieving the value associated with key x in finite mapping h and $h(x) = \perp$ denotes that no value is associated with x . Notation $h \uplus h'$ denotes the right-biased union of two finite mappings, i.e. when both maps have the same key x , we keep the value $h'(x)$. Following common practice, notation $h[x \mapsto v]$ denotes finite mapping update, i.e. finite mapping h' such that: 1) $h'(x) = v$ and 2) $h'(y) = h(y)$, for $x \neq y$.

The proposed semantics is based on Transactional Locking 2 algorithm (TL2) (Dice *et al.*, 2006). We define it using a small-step semantics defined in Figures 5 and 7. We use evaluation contexts, which are defined in Figure 2, to avoid the need of “congruence” rules in semantics.

Evaluation contexts for transactions

$\mathbb{T}[\cdot] ::=$	$\text{writeTVar } x \ \mathbb{T}[\cdot]$
	$\mathbb{T}[\cdot] \oplus t$
	$v \oplus \mathbb{T}[\cdot]$
	$\text{if } \mathbb{T}[\cdot] \text{ then } t \text{ else } t'$

Evaluation contexts for processes

$\mathbb{P}[\cdot] ::=$	$\text{fork } \mathbb{P}[\cdot]$
	$\mathbb{P}[\cdot] \oplus t$
	$t \oplus \mathbb{P}[\cdot]$

Fig. 2. Evaluation contexts for high-level language.

Informally, TL2 algorithm works as follows: threads execute reads and writes to objects, but no memory locations are actually modified. All writes and reads are recorded in write and read logs. When a transaction finishes, it validates its log to check if it has seen a consistent view of memory, and its changes are committed to memory.

$$\Theta(x, i) = \begin{cases} (v, \Theta) & \text{if } \Theta_w(x) = v \\ (v, \Theta) & \text{if } \Theta_w(x) = \perp, \Theta_r(x) \neq \perp, \Theta_h(x) = (i', v) \text{ and } i \geq i' \\ (v, \Theta_r[x \mapsto v]) & \text{if } \Theta_w(x) = \Theta_r(x) = \perp, \Theta_h(x) = (i', v), \text{ and } i \geq i' \\ \text{retry} & \text{if } \Theta_h(x) = (i', v) \text{ and } i < i' \end{cases}$$

Fig. 3. Reading a variable

$$\text{consistent}(\Theta, i) = \forall x. \Theta_r(x) = (i, v) \rightarrow \Theta_h(x) = (i', v) \wedge i \geq i'$$

Fig. 4. Predicate for consistency of transaction logs

Verifying consistency of a transaction.

Now the small step semantics for processes

References

- Claessen, Koen, & Hughes, John. (2000). Quickcheck: A lightweight tool for random testing of haskell programs. *Pages 268–279 of: Proceedings of the fifth acm sigplan international conference on functional programming*. ICFP '00. New York, NY, USA: ACM.
- Dice, Dave, Shalev, Ori, & Shavit, Nir. (2006). Transactional locking ii. *Pages 194–208 of: Proceedings of the 20th international conference on distributed computing*. DISC'06. Berlin, Heidelberg: Springer-Verlag.
- Gill, Andy, & Runciman, Colin. (2007). Haskell program coverage. *Pages 1–12 of: Proceedings of the acm sigplan workshop on haskell workshop*. Haskell '07. New York, NY, USA: ACM.
- Harris, Tim, Marlow, Simon, Peyton-Jones, Simon, & Herlihy, Maurice. (2005). Composable memory transactions. *Pages 48–60 of: Proceedings of the tenth acm sigplan symposium on principles and practice of parallel programming*. PPOPP '05. New York, NY, USA: ACM.
- Hu, Liyang, & Hutton, Graham. (2009). Towards a Verified Implementation of Software Transactional Memory. Achten, Peter, Koopman, Pieter, & Morazan, Marco (eds), *Trends in Functional Programming volume 9*. Intellect. Selected papers from the Ninth Symposium on Trends in Functional Programming, Nijmegen, The Netherlands, May 2008.
- Shavit, Nir, & Touitou, Dan. (1995). Software transactional memory. *Pages 204–213 of: Proceedings of the fourteenth annual acm symposium on principles of distributed computing*. PODC '95. New York, NY, USA: ACM.

$$\boxed{\langle \Theta, \sigma, t \rangle \mapsto_{T_i} \langle \Theta', \sigma', t' \rangle}$$

$$\begin{array}{c}
\frac{\langle \Theta, \sigma, t \rangle \mapsto_{T_i} \langle \Theta', \sigma', t' \rangle}{\langle \Theta, \sigma, \mathbb{T}[t] \rangle \mapsto_{T_i} \langle \Theta', \sigma', \mathbb{T}[t'] \rangle} \text{ (Context)} \qquad \frac{\sigma' = \text{Write } i \ x \ v :: \sigma}{\langle \Theta, \sigma, \text{writeTVar } x \ v \rangle \mapsto_{T_i} \langle \Theta_w[x \mapsto v], \sigma', t' \rangle} \text{ (Write)} \\
\\
\frac{\sigma' = \text{Abort } i :: \sigma}{\langle \Theta, \sigma, \text{writeTVar } x \ \text{retry} \rangle \mapsto_{T_i} \langle \Theta, \sigma', \text{retry} \rangle} \text{ (WriteRetry)} \qquad \frac{\sigma' = \text{Read } i \ x \ v :: \sigma \quad (v, \Theta') = \Theta(x, i)}{\langle \Theta, \sigma, \text{readTVar } x \rangle \mapsto_{T_i} \langle \Theta', \sigma', v \rangle} \text{ (ReadVal)} \\
\\
\frac{\sigma' = \text{Abort } i :: \sigma \quad \text{retry} = \Theta(x, i)}{\langle \Theta, \sigma, \text{readTVar } x \rangle \mapsto_{T_i} \langle \Theta, \sigma', \text{retry} \rangle} \text{ (ReadRetry)} \qquad \frac{}{\langle \Theta, \sigma, v \oplus v' \rangle \mapsto_{T_i} \langle \Theta, \sigma, v + v' \rangle} \text{ (Plus)} \\
\\
\frac{\langle \Theta, \sigma, t \rangle \mapsto_{T_i}^* \langle \Theta', \sigma', v \rangle}{\langle \Theta, \sigma, t \ \text{orElse } t' \rangle \mapsto_{T_i} \langle \Theta', \sigma', v \rangle} \text{ (OrElseVal)} \qquad \frac{\langle \Theta, \sigma, t \rangle \mapsto_{T_i}^* \langle \Theta', \sigma', \text{retry} \rangle}{\langle \Theta, \sigma, t \ \text{orElse } t' \rangle \mapsto_{T_i} \langle \Theta', \sigma', \text{retry} \rangle} \text{ (OrElseRetry)} \\
\\
\frac{\sigma' = \text{Abort } i :: \sigma}{\langle \Theta, \sigma, \text{retry} \oplus t \rangle \mapsto_{T_i} \langle \Theta, \sigma', \text{retry} \rangle} \text{ (PlusRetryL)} \qquad \frac{\sigma' = \text{Abort } i :: \sigma}{\langle \Theta, \sigma, t \oplus \text{retry} \rangle \mapsto_{T_i} \langle \Theta, \sigma', \text{retry} \rangle} \text{ (PlusRetryR)} \\
\\
\frac{}{\langle \Theta, \sigma, \text{if } v = 0 \text{ then } t \text{ else } t' \rangle \mapsto_{T_i} \langle \Theta, \sigma, t \rangle} \text{ (IfZero)} \\
\frac{}{\langle \Theta, \sigma, \text{if } v \neq 0 \text{ then } t \text{ else } t' \rangle \mapsto_{T_i} \langle \Theta, \sigma, t' \rangle} \text{ (IfNonZero)} \\
\frac{\sigma' = \text{Abort } i :: \sigma}{\langle \Theta, \sigma, \text{if } \text{retry} \text{ then } t \text{ else } t' \rangle \mapsto_{T_i} \langle \Theta, \sigma', \text{retry} \rangle} \text{ (IfRetry)}
\end{array}$$

Fig. 5. Small-step operational semantics for transactions.

$$\boxed{\langle \Theta, \sigma, t \rangle \mapsto_{T_i}^* \langle \Theta', \sigma', t' \rangle}$$

$$\begin{array}{c}
\frac{}{\langle \Theta, \sigma, t \rangle \mapsto_{T_i}^* \langle \Theta, \sigma, t \rangle} \text{ (TRefI)} \\
\\
\frac{\langle \Theta, \sigma, t \rangle \mapsto_{T_i} \langle \Theta_1, \sigma_1, t_1 \rangle \quad \langle \Theta_1, \sigma_1, t_1 \rangle \mapsto_{T_i}^* \langle \Theta', \sigma', t' \rangle}{\langle \Theta, \sigma, t \rangle \mapsto_{T_i}^* \langle \Theta', \sigma', t' \rangle} \text{ (TTran)}
\end{array}$$

Fig. 6. Reflexive-transitive closure of transaction small step semantics.

$$\boxed{\langle h, \sigma, i, s \rangle \mapsto_P \langle h', \sigma', i', s' \rangle}$$

$$\frac{\langle h, \sigma, i, s \rangle \mapsto_P \langle h', \sigma', i', s' \rangle}{\langle h, \sigma, i, p : s \rangle \mapsto_P \langle h', \sigma', i', p : s' \rangle} \text{ (Preempt)} \qquad \frac{\langle h, \sigma, i, p : s \rangle \mapsto_P \langle h', \sigma', i', p' : s' \rangle}{\langle h, \sigma, i, \mathbb{P}[p] : s \rangle \mapsto_P \langle h', \sigma', i, \mathbb{P}[p'] : s \rangle} \text{ (Context)}$$

$$\frac{}{\langle h, \sigma, i, v \oplus v' : s \rangle \mapsto_P \langle h, \sigma, i, v + v' : s \rangle} \text{ (Plus)} \qquad \frac{}{\langle h, \sigma, i, \text{fork } p : s \rangle \mapsto_P \langle h, \sigma, i, 0 : p : s \rangle} \text{ (Fork)}$$

$$\frac{\text{consistent}(\Theta', i) \quad h' = h \uplus \Theta_w \quad \langle (h, \bullet, \bullet), \sigma, t \rangle \mapsto_{T_i}^* \langle \Theta', \sigma', v \rangle}{\langle h, \sigma, i, \text{Atomic } t : s \rangle \mapsto_P \langle h', \sigma', i + 1, v : s \rangle} \text{ (AtomicVal)}$$

$$\frac{\langle (h, \bullet, \bullet), \sigma, t \rangle \mapsto_{T_i}^* \langle \Theta', \sigma', \text{retry} \rangle}{\langle h, \sigma, i, \text{Atomic } t : s \rangle \mapsto_P \langle \sigma', i, h, 0 : s \rangle} \text{ (AtomicRetry)} \qquad \frac{\langle (h, \bullet, \bullet), \sigma, t \rangle \mapsto_{T_i}^* \langle \Theta', \sigma', t' \rangle \quad \neg \text{consistent}(\Theta', i)}{\langle h, \sigma, i, \text{Atomic } t : s \rangle \mapsto_P \langle \sigma', i, h, 0 : s \rangle} \text{ (AtomicNon)}$$

Fig. 7. Small-step operational semantics for processes.

