

An Opaque Model for Software Transactional Memory for Haskell

Rodrigo Ribeiro ¹ André Du Bois ²

DECOM, Universidade Federal de Ouro Preto (UFOP), Ouro Preto

CDTec, Universidade Federal de Pelotas(UFPel), Pelotas

2 de maio de 2017

Summary

Brief Bio

Introduction

Motivation

Software Transactional Memory in Haskell

Transactional Memory Correctness

Formal semantics for transactional memory

Core language definition

Definitions and notations used

TL2 semantics

Checking Opacity

Conclusions

Brief Bio — (I)

- ▶ MSc in Computer Science (UFMG) - 2007;
- ▶ DSc in Computer Science (UFMG) - 2013;
- ▶ Lecturer at Universidade Federal de Ouro Preto.
- ▶ Post-doc (UFPel) - 2016 - present .

Brief Bio — (II)

- ▶ Main research topic: Type theory.
- ▶ Applications:
 - ▶ Programming languages design
 - ▶ Formal verification — proof assistants.

Currently working in... — (III)

- ▶ Formal semantics for STM.
- ▶ Formal verification of Featherweight Java typechecker.
- ▶ Formal semantics and verification of DSLs for temporal media (music, video, etc).
- ▶ Other projects
 - ▶ Formal verification of parsing algorithms.
 - ▶ Compilation of partial C programs.

This Talk — (IV)

- ▶ Formal semantics for STM.
 - ▶ Needed to ensure safety properties of STM.
 - ▶ Results reported in a paper submitted to Haskell Symposium 2017.
 - ▶ We assume that audience is familiar with Haskell syntax and monadic programming.

Multicores are coming! — (V)

- ▶ For 50 years, hardware designers delivered 40-50% increases per year in sequential program performance.
- ▶ Around 2004, this promise fails because power and cooling issues made impossible to increase clock frequencies.
- ▶ Result: If we want better performance, parallelism is no longer an option!

Parallelism and multicores — (VI)

- ▶ Parallelism is essential to improve performance on a multi-core machine.
- ▶ Unfortunately, parallel programming is immensely more error-prone than traditional sequential programming...
- ▶ How to do it? Locks and conditional variables.

Locks and programming — (VII)

- ▶ **Correct** use of locks can solve concurrency problems, but...
 - ▶ Locks are amazingly hard to use correctly!
 - ▶ Incorrect use of locks can cause races and deadlocks which are difficult to debug.
- ▶ Better solution: Transactions!

Transactional Memory — (VIII)

- ▶ Intuitively, write sequential code and wrap it using an **atomic** block.
- ▶ Atomic blocks execute in isolation.
 - ▶ with automatic retry if another conflicting atomic block interferes.

How does it work? — (IX)

- ▶ The TM runtime tries to interleave transactions.
- ▶ How? Here's one way:
 - ▶ Read and writes happens in a transaction local log.
 - ▶ At the end, transactions validate its log.
 - ▶ If validation succeeds, changes are committed to memory
 - ▶ If fails, re-runs from the beginning, discarding changes.

Our focus: TL2 algorithm — (X)

- ▶ Uses logs and a global clock.
- ▶ Every transaction holds its read stamp.
 - ▶ Global clock value when transaction started.
- ▶ Every transactional variable has a write stamp.
 - ▶ When reading, if write stamp $>$ read stamp - abort.

Why Haskell? — (XI)

- ▶ Haskell is pure functional language, i.e., by default, side-effects are available only through monads.
- ▶ This allows for a strict separation between effectful and pure code.
- ▶ Using monads we can isolate transactional code from other effects present in program.

Concurrency in Haskell — (XII)

- ▶ Function `fork` spawns a new thread.
- ▶ It takes an action as its argument.

`fork :: IO a → IO ThreadId`

```
main = do {  
    fork someAction;  
    anotherAction;  
    ...}
```

Atomic blocks in Haskell — (XIII)

- ▶ Idea: use a function `atomic` that guarantees atomic execution of its argument computation atomically.

```
main = do {  
    r ← new 0;  
    fork (atomic (someAction));  
    atomic (anotherAction); ... }
```

STM Haskell interface — (XIV)

- ▶ Transactional variables:
 - ▶ **data** TVar *a* = ...
- ▶ Transactional memory monad:
 - ▶ **data** STM *a* = ...
- ▶ Creating/ reading / writing variables.

`newTVar` :: *a* → STM (TVar *a*)

`readTVar` :: TVar *a* → STM *a*

`writeTVar` :: TVar *a* → *a* → STM ()

STM Haskell interface — (XV)

- ▶ User controlled abort of transactions.

`retry` :: STM *a*

- ▶ Choice operator.

`orElse` :: STM *a* → STM *a* → STM *a*

- ▶ Running a transaction.

`atomically` :: STM *a* → IO *a*

STM Haskell example — (XVI)

```
type Var = TVar Float
```

```
transferMoney :: Float → Var → Var → STM ()
```

```
transferMoney amount acc1 acc2
```

```
  = do
```

```
    v ← readTVar acc1
```

```
  if v ≥ amount
```

```
    then do
```

```
      writeTVar acc1 (v − amount)
```

```
      v2 ← readTVar acc2
```

```
      writeTVar acc2 (v2 + amount)
```

```
    else retry
```

Transactional memory summary — (XVII)

- ▶ TM provides a simple programming model to write concurrent code.
- ▶ STM simplicity (for programmer) has a price:
 - ▶ Designing efficient algorithms for TM is an art.
 - ▶ Implementations usually use subtle, but efficient, algorithms.
 - ▶ How can we guarantee safety of an TM implementation?
 - ▶ What means “correctness” for TM?

Correctness criteria for STM — (XVIII)

- ▶ Several criteria proposed in literature.
- ▶ They are based on the concept of TM **history**.
- ▶ History:
 - ▶ Sequence of operations executed during a TM run.
- ▶ Example:

$$H = \left| \begin{array}{l} T_1.read\ x \rightarrow 0, \\ T_1.write(x, 1), \\ T_2.read(x) \rightarrow 1 \\ T_1.commit \\ T_2.abort \end{array} \right|$$

Correctness criteria for STM — (XIX)

- ▶ How histories can be used to establish correctness?
- ▶ A correct history should be such that:
 - ▶ Committed transactions should appear to be executed with no other transactions.
 - ▶ Aborted transactions should leave no effect.
 - ▶ Possible to find a history H' s.t. there's a **total order** in which transactions appear to be executed.

Correctness criteria for STM — (XX)

- ▶ How histories can be used to establish correctness?
- ▶ A correct history should be such that:
 - ▶ Possible to find a history H' s.t. there's a **total order** in which transactions appear to be executed.
- ▶ Here we'll focus on a criteria named **opacity**.
- ▶ A TM algorithm is correct if all histories generated by it are opaque.

A glimpse of our proposal — (XXI)

- ▶ Define a TM algorithm as a small-step operational semantics over a core language.
- ▶ Semantics should produce the TM history.
- ▶ Verify that all produced histories satisfy the correctness criteria.
 - ▶ We use property based testing for this.

Core language — (XXII)

- values

newtype Val = Val { unVal :: Int }

- variables

newtype Var = Var { unVar :: Int }

- transaction identifiers

newtype Id = Id { unId :: Int }

- clock values

newtype Stamp = Stamp { unStamp :: Int }

Core language — (XXIII)

data Tran =

| | |
|----------------------|-----------------|
| TVal Val | - values |
| TRead Var | - read op. |
| TWrite Var Tran | - write op. |
| Tran \oplus_T Tran | - composition |
| TIf Tran Tran Tran | - conditional |
| TOrElse Tran Tran | - orElse op. |
| TRetry | - user abort |
| TAbort | - runtime abort |

Core language — (XXIV)

- transaction id + global clock info.

type TIdent = Maybe (Id, Stamp)

data Proc =

| | | | |
|------|----------|-----------------|----------------|
| PVal | Val | - values | |
| | PFork | Proc | - forking |
| | PAAtomic | TIdent Tran | - atomic block |
| | Proc | \oplus_P Proc | - composition |

Transactional histories — (XXV)

```
data Event =  
  IRead Id Var Val - reading a value  
  | IWrite Id Var Val - writing a value  
  | IBegin Id      - beginning a transaction.  
  | ICommit Id     - commit a transaction.  
  | IAbort Id      - runtime abort.  
  | IRetry Id      - user abort.  
type History = [Event]
```

Notations: Finite maps — (XXVI)

- ▶ Empty finite map \bullet .
- ▶ Lookup:

$$h(x) = \begin{cases} v & \text{if } [x \mapsto v] \in h \\ \perp & \text{otherwise} \end{cases}$$

- ▶ Right biased union: $h \uplus h'$.
- ▶ Removing: $h \mid_x = h - [x \mapsto v]$, for some v .

Notations: TM state — (XXVII)

- ▶ Heaps, read and write sets represented as finite maps.
- ▶ $\Theta = \langle \Theta_h, \Theta_r, \Theta_w, \Theta_T \rangle$, where:
 - ▶ Θ_h : heap
 - ▶ Θ_r and Θ_w : finite maps between trans. id and its read / write sets.
 - ▶ Θ_T : finite map between trans. id and the transaction itself.

Reading a variable: intuition — (XXVIII)

- ▶ $\Theta(x, i, j)$: reads the value of variable x in transaction j which has read stamp i .
- ▶ How reading proceeds:
 - ▶ If x is in j write set, return its value.
 - ▶ If x is in j read set and i isn't less than x write stamp, return x associated value.
 - ▶ Lookup value in heap. If x value has a stamp greater than i , abort.
 - ▶ Otherwise, update read-set.

Reading a variable, formally — (XXIX)

$$\Theta(x, i, j) = \left\{ \begin{array}{ll} (v, \Theta) & \text{if } \Theta_w(x, j) = (i', v) \\ (v, \Theta) & \text{if } \Theta_w(x, j) = \perp, \\ & \Theta_r(x, j) = (i', v), i \geq i' \\ (v, \Theta_r[j, x \mapsto v]) & \text{if } \Theta_w(x) = \Theta_r(x) = \perp, \\ & \Theta_h(x) = (i', v), \text{ and } i \geq i' \\ \text{TAbort} & \text{if } \Theta_h(x) = (i', v) \text{ and } i < i' \end{array} \right.$$

Consistency check — (XXX)

- ▶ *consistent*(Θ, i, j) holds if transaction j has finished its execution on a valid view of memory.

$$\text{consistent}(\Theta, i, j) = \forall x. \Theta_r(x, j) = (i, v) \rightarrow \Theta_h(x) = (i', v) \wedge i \geq i'$$

Semantics — (XXXI)

- ▶ Transaction semantics uses triples $\langle \Theta, \sigma, t \rangle$:
 - ▶ Θ : heap and TM logs.
 - ▶ σ : TM execution history
 - ▶ t : transaction being executed.
 - ▶ i : current read stamp
 - ▶ j : current transaction id.

$$\langle \Theta, \sigma, t \rangle \mapsto_{T_{ij}} \langle \Theta', \sigma', t' \rangle$$

Semantics — (XXXII)

- ▶ Process semantics uses 5-uples

- ▶ Θ : heap and TM logs.
- ▶ σ : TM execution history
- ▶ i : global clock
- ▶ j : last transaction id used.
- ▶ p : process being executed.

$$\langle \Theta, \sigma, j, i, t \rangle \mapsto_P \langle \Theta', \sigma', j', i', t' \rangle$$

Evaluation Contexts — (XXXIII)

$$\begin{array}{lcl} \mathbb{T}[\cdot] & ::= & \text{TWrite } v \ \mathbb{T}[\cdot] \\ & | & \mathbb{T}[\cdot] \oplus_{\mathbb{T}} t \\ & | & \text{TVal } v \oplus_{\mathbb{T}} \mathbb{T}[\cdot] \\ & | & \text{TIf } \mathbb{T}[\cdot] \ t \ t' \\ & | & \text{TOrElse } \mathbb{T}[\cdot] \ t \end{array}$$
$$\begin{array}{lcl} \mathbb{P}[\cdot] & ::= & \text{PFork } \mathbb{P}[\cdot] \\ & | & \mathbb{P}[\cdot] \oplus_{\mathbb{P}} t \\ & | & \text{PVal } v \oplus_{\mathbb{P}} \mathbb{P}[\cdot] \\ & | & \text{PAtomic } \mathbb{P}[\cdot] \end{array}$$

Evaluation context reduction — (XXXIV)

$$\frac{\langle \Theta, \sigma, t \rangle \mapsto_{T_{ij}} \langle \Theta', \sigma', t' \rangle}{\langle \Theta, \sigma, \mathbb{T}[t] \rangle \mapsto_{T_{ij}} \langle \Theta', \sigma', \mathbb{T}[t'] \rangle} \quad (TContext)$$

$$\frac{\langle \Theta, \sigma, j, i, t \rangle \mapsto_P \langle \Theta', \sigma', j', i', t' \rangle}{\langle \Theta, \sigma, j, i, \mathbb{P}[p] : s \rangle \mapsto_P \langle \Theta, \sigma', j', i', \mathbb{P}[p'] : s \rangle} \quad (PContext)$$

Read Semantics — (XXXV)

$$\frac{\Theta(v, i, j) = (val, \Theta') \quad \sigma' = \text{IRead } j \ v \ val : \sigma}{\langle \Theta, \sigma, \text{TRead } v \rangle \mapsto_{T_{ij}} \langle \Theta', \sigma', \text{TVal } val \rangle} \quad (TReadOk)$$

$$\frac{\Theta(v, i, j) = (\text{TAabort}, \Theta') \quad \sigma' = \text{IAabort } j : \sigma}{\langle \Theta, \sigma, \text{TRead } v \rangle \mapsto_{T_{ij}} \langle \Theta, \sigma', \text{TAabort} \rangle} \quad (TReadFail)$$

Write Semantics — (XXXVI)

$$\frac{\begin{array}{l} t = \text{TVal } val \\ \sigma' = \text{IWrite } j \ v \ val : \sigma \\ \Theta' = \langle \Theta_h, \Theta_r, \Theta_w[j, x \mapsto val], \Theta_T \rangle \end{array}}{\langle \Theta, \sigma, \text{TWrite } v \ t \rangle \mapsto_{T_{ij}} \langle \Theta', \sigma', t \rangle} \quad (\text{TWriteVal})$$

$$\frac{}{\langle \Theta, \sigma, \text{TWrite } v \ \text{TFail} \rangle \mapsto_{T_{ij}} \langle \Theta, \sigma, \text{TFail} \rangle} \quad (\text{TWriteFail})$$

Composition Semantics — (XXXVII)

$$\frac{val = val_1 + val_2}{\langle \Theta, \sigma, (\mathbf{TVal} \ val_1) \oplus_{\mathbf{T}} (\mathbf{TVal} \ val_2) \rangle \mapsto_{T_{ij}} \langle \Theta, \sigma, \mathbf{TVal} \ val \rangle} \quad (TAddVal)$$

$$\frac{}{\langle \Theta, \sigma, \mathbf{TFail} \oplus_{\mathbf{T}} t \rangle \mapsto_{T_{ij}} \langle \Theta', \sigma', \mathbf{TFail} \rangle} \quad (TAddL)$$

$$\frac{}{\langle \Theta, \sigma, (\mathbf{TVal} \ val) \oplus_{\mathbf{T}} \mathbf{TFail} \rangle \mapsto_{T_{ij}} \langle \Theta', \sigma', \mathbf{TFail} \rangle} \quad (TAddR)$$

Conditional Semantics — (XXXVIII)

$$\frac{}{\langle \Theta, \sigma, \text{TIf}(\text{TVal } 0) \, t \, t' \rangle \mapsto_{T_{ij}} \langle \Theta, \sigma, t \rangle} \text{ (TIfZero)}$$

$$\frac{v \neq 0}{\langle \Theta, \sigma, \text{TIf}(\text{TVal } v) \, t \, t' \rangle \mapsto_{T_{ij}} \langle \Theta, \sigma, t' \rangle} \text{ (TIfNonZero)}$$

$$\frac{}{\langle \Theta, \sigma, \text{TIf TFail} \, t \, t' \rangle \mapsto_{T_{ij}} \langle \Theta, \sigma, \text{TFail} \rangle} \text{ (TIfFail)}$$

OrElse Semantics — (XXXIX)

$$\frac{}{\langle \Theta, \sigma, \text{TOrElse}(\text{TVal } v) \ t' \rangle \mapsto_{T_{ij}} \langle \Theta', \sigma', \text{TVal } v \rangle} \text{ (TOrElseVal)}$$

$$\frac{}{\langle \Theta, \sigma, \text{TOrElse} \text{TRetry } t' \rangle \mapsto_{T_{ij}} \langle \Theta, \sigma, t' \rangle} \text{ (TOrElseR)}$$

$$\frac{}{\langle \Theta, \sigma, \text{TOrElse} \text{TAabort } t' \rangle \mapsto_{T_{ij}} \langle \Theta, \sigma, \text{TAabort} \rangle} \text{ (TOrElseA)}$$

Process preemption — (XL)

$$\frac{\langle \Theta, \sigma, j, i, s \rangle \mapsto_P \langle \Theta', \sigma', j', i', s' \rangle}{\langle \Theta, \sigma, j, i, p : s \rangle \mapsto_P \langle \Theta', \sigma', j', i', p : s' \rangle} \text{ (PPreempt)}$$

Forking semantics — (XLI)

$$\frac{s' = \text{PVal } 0 : p : s}{\langle \Theta, \sigma, j, i, (\text{PFork } p) : s \rangle \mapsto_P \langle \Theta, \sigma, j, i, s' \rangle} \text{ (PFork)}$$

Atomic block semantics — (XLII)

$$\sigma' = \text{IBegin } j : \sigma$$

$$s' = s \# [\text{PAtomic } (i, j) \ t]$$

$$\frac{\Theta_1 = \langle \Theta_h, \Theta_r [j \mapsto \bullet], \Theta_w [j \mapsto \bullet], \Theta_T [j \mapsto t] \rangle}{\langle \Theta, \sigma, j, i, \text{PAtomic } () \ t : s \rangle \mapsto_P \langle \Theta', \sigma', j+1, i, s' \rangle} \text{ (PAt1)}$$

Atomic block semantics — (XLIII)

$$\frac{\begin{array}{c} v = \text{TVa}l\ n \qquad \text{consistent}(\Theta, i, j) \\ \sigma' = \text{ICommit}\ j:\sigma \quad \Theta' = \langle \Theta'_h, \Theta_r|_j, \Theta_w|_j \rangle \\ \Theta'_h = \Theta_h \uplus \Theta_w(j) \end{array}}{\langle \Theta, \sigma, j, i, \text{PAtomic}\ v:s \rangle \mapsto_P \langle \Theta', \sigma', j, i+1, \text{PVal}\ n:s \rangle} \text{ (PA}t2\text{)}$$

Atomic block semantics — (XLIV)

$$\frac{\begin{array}{l} \Theta' = \langle \Theta_h, \Theta_r|_j, \Theta_w|_j, \Theta_T|_j \rangle \\ s' = s \mathbin{++} \text{PAtomic}() \ t \quad t = \Theta_t(j) \end{array}}{\langle \Theta, \sigma, j, i, \text{PAtomic TFail}:s \rangle \mapsto_P \langle \Theta', \sigma, j+1, i, s' \rangle} \text{ (PAt3)}$$

Checking Opacity — (XLV)

- ▶ We check safety properties using Quickcheck, a Haskell library for property based testing.
- ▶ Library formed by combinators for building generators and properties.
 - ▶ Generators build random values which are checked against relevant properties.
 - ▶ Combinators for properties mimics first-order logic connectives and quantifiers.

Defining Opacity — (XLVI)

- ▶ A TM algorithm is opaque if all prefixes of its generated histories are final state opaque.
- ▶ In Haskell:

```
opacity :: History → Bool  
opacity = all finalStateOpacity ∘ inits
```


Defining Opacity — (XLVII)

- ▶ A history is final state opaque if exists some completion of it that preserves real time order and all of its transactions are legal.
- ▶ In Haskell:

`finalStateOpacity :: History → Bool`

`finalStateOpacity`

`= some prop ∘ completions`

where

`prop tr = preservesOrder tr ∧ legal tr`

`some p xs = (null xs) ∨ (any p xs)`

Defining Opacity — (XLVIII)

- ▶ Completion of a history H is a history S , s.t.
 - ▶ All live and non-commit pending transactions of H are aborted in S ; and
 - ▶ All commit pending transactions of H are aborted or committed in S .
- ▶ Our model we do not consider commit-pending transactions.
- ▶ Complete a history is just abort all live transactions.

Defining Opacity — (XLIX)

- ▶ Real time order: $t_k \prec_H t_m$
 - ▶ t_k is completed and its last event occurs before t_m s first.
- ▶ A history H' preserves the real time ordering of H :

$$\forall t_k t_m. t_k \prec_H t_m \rightarrow t_k \prec_{H'} t_m$$

Defining Opacity — (L)

- ▶ Legal histories.
 - ▶ A history is said to be *legal* if all values read after a write in a variable are equal to last value written

Checking Opacity — (LI)

- ▶ We use function `opacity` as a Quickcheck property over histories produced by executing random generated programs.
- ▶ Semantics passes in all test-suite runs.

Checking Opacity — (LII)

- ▶ So, our semantics does enjoy opacity?
 - ▶ As pointed by Dijkstra: “Tests can only prove the presence of bugs...”.
- ▶ In order to affirm this categorically, we need a formal proof.
 - ▶ A venue that we intend to pursue in future work.

History not told... (LIII)

- ▶ Alternative STM-Haskell based semantics (does not have opacity).
- ▶ Compilation to virtual machine.
- ▶ Another correctness criteria: markability.
- ▶ Other properties verified
 - ▶ Compilation preserves safety properties.
 - ▶ Equivalence of safety properties.

Conclusions — (LIV)

- ▶ We define small-step operational semantics for STM-Haskell like language.
 - ▶ TL2-based semantics.
 - ▶ STM-Haskell based semantics.
- ▶ Implemented interpreters for the semantics and relevant properties over it using Haskell.

Conclusions — (LV)

- ▶ Property based testing.
 - ▶ Allows for quickly identify errors in semantics — counter-examples.
 - ▶ Provides some degree of assurance, before using more formal approaches.